

# Task-agnostic Amortized Multi-Objective Optimization

Xinyu Zhang, Conor Hassan, Julien Martinelli, Daolang Huang, Samuel Kaski

ICLR 2026 submission ↗

October 30<sup>th</sup>, 2025

"Amortize everything 😱"

— Daolang Huang

## Black-box optimization

Let  $\mathcal{X}$  be a design space and  $f$  an objective function, we seek

$$x = \operatorname{argmax}_{x \in \mathcal{X}} f(x)$$

## Black-box optimization

Let  $\mathcal{X}$  be a design space and  $f$  an objective function, we seek

$$\mathbf{x} = \operatorname{argmax}_{\mathbf{x} \in \mathcal{X}} f(\mathbf{x})$$

- No information about  $f$ : gradients, convexity, invariances...
- Can only evaluate  $y_i = f(\mathbf{x}_i) + \varepsilon$
- Collecting observations  $\mathcal{D}_t = \{(\mathbf{x}_i, y_i)\}_{i=1}^t$

## Black-box optimization

Let  $\mathcal{X}$  be a design space and  $f$  an objective function, we seek

$$x = \operatorname{argmax}_{x \in \mathcal{X}} f(x)$$

- No information about  $f$ : gradients, convexity, invariances...
- Can only evaluate  $y_i = f(x_i) + \varepsilon$
- Collecting observations  $\mathcal{D}_t = \{(x_i, y_i)\}_{i=1}^t$

Bayesian Optimization tackles this problem by learning a *cheap-to-evaluate statistical surrogate* of  $f$

$$f(x) \sim \mathcal{GP}(\mu_{\theta_1}(x), k_{\theta_2}(x, x'))$$

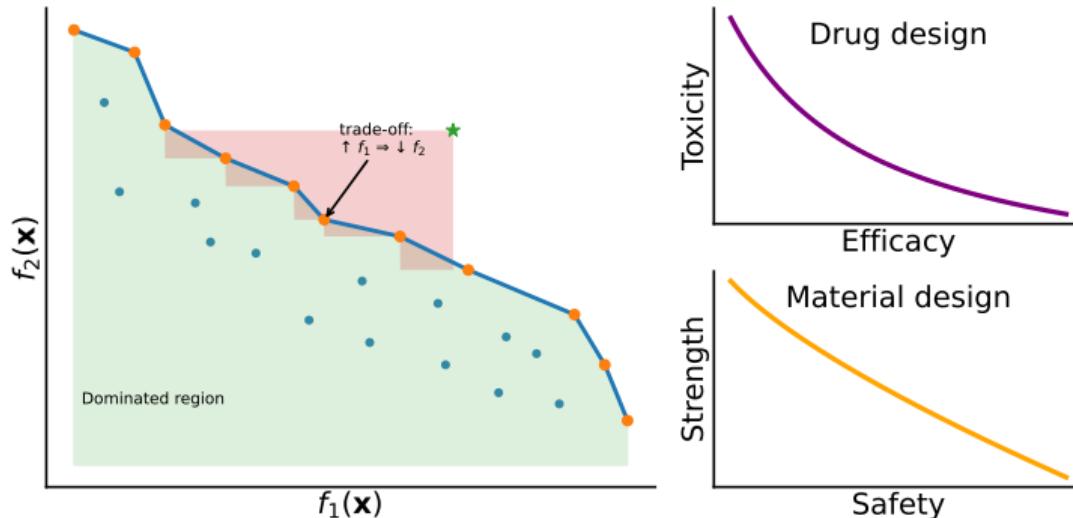
Yields next sample to query  $x_t$  selected sequentially based on an **acquisition function**

$$x_{t+1} = \operatorname{argmax}_{x \in \mathcal{X}} \alpha_{\theta_3}(x | \mathcal{D}_t)$$

Evaluate  $f(x_{t+1})$ ; append  $\mathcal{D}_{t+1} \leftarrow \mathcal{D}_t \cup \{(x_{t+1}, y_{t+1})\}$ ; update surrogate  $p(f | \mathcal{D}_{t+1})$ , repeat until satisfied.

# Multi-objective optimization

We now observe a vector of objectives  $f(\mathbf{x}) = [f_1(\mathbf{x}), \dots, f_M(\mathbf{x})] \in \mathbb{R}^M$



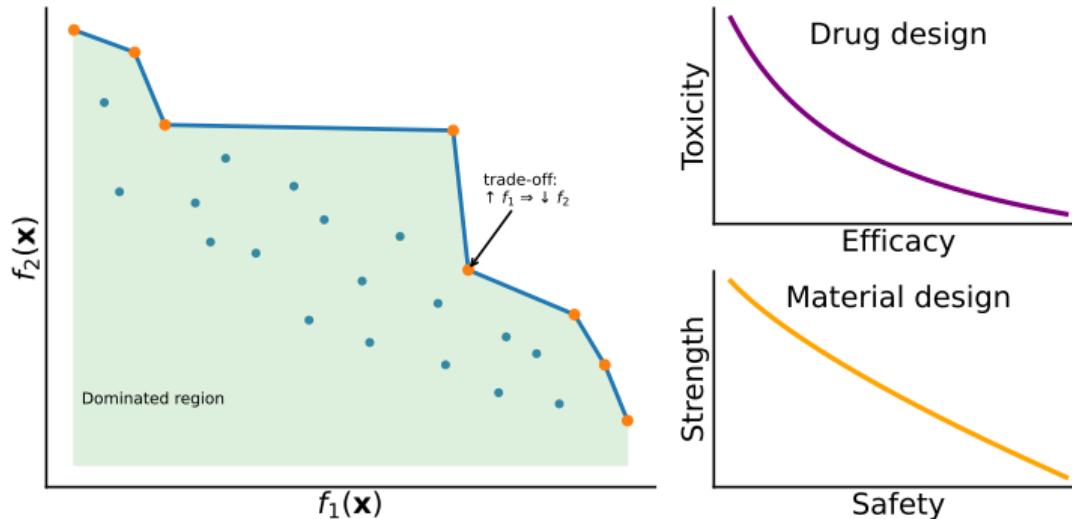
$\mathcal{P} \subset \mathbb{R}^m$  set of objective vectors and reference  $\mathbf{r} \in \mathbb{R}^m$  with  $\mathbf{r} \leq \mathbf{p}$  for all  $\mathbf{p} \in \mathcal{P}$  (e.g.,  $r_i \leq \min_{\mathbf{p} \in \mathcal{P}} p_i$ ).

$$HV(\mathcal{P}; \mathbf{r}) = \lambda_m \left( \bigcup_{\mathbf{p} \in \mathcal{P}} \prod_{i=1}^m [r_i, p_i] \right)$$

$$HVI(\mathbf{x} \mid \mathcal{P}, \mathbf{r}) = HV(\mathcal{P} \cup \{\mathbf{x}\}; \mathbf{r}) - HV(\mathcal{P}; \mathbf{r})$$

# Multi-objective optimization

We now observe a vector of objectives  $\mathbf{f}(\mathbf{x}) = [f_1(\mathbf{x}), \dots, f_M(\mathbf{x})] \in \mathbb{R}^M$



$\mathcal{P} \subset \mathbb{R}^m$  set of objective vectors and reference  $\mathbf{r} \in \mathbb{R}^m$  with  $\mathbf{r} \leq \mathbf{p}$  for all  $\mathbf{p} \in \mathcal{P}$  (e.g.,  $r_i \leq \min_{\mathbf{p} \in \mathcal{P}} p_i$ ).

$$HV(\mathcal{P}; \mathbf{r}) = \lambda_m \left( \bigcup_{\mathbf{p} \in \mathcal{P}} \prod_{i=1}^m [r_i, p_i] \right)$$

$$HVI(\mathbf{x} \mid \mathcal{P}, \mathbf{r}) = HV(\mathcal{P} \cup \{\mathbf{x}\}; \mathbf{r}) - HV(\mathcal{P}; \mathbf{r})$$

## Caveats

- ① Performances highly dependent on surrogate and acquisition function pair
  - ▶ Requires careful, expert selection of the ideal combination
  - ▶ Will not transfer to the next problem

→ "Learn" the language of optimization
- ② Slow in high-throughput settings due to vanilla GP cubic complexity

→ Do a huge offline pre-training step, reduce inference to a feedforward pass
- ③ Most of the time myopic, focused on 1-step optimality

→ Train over long optimization horizons using RL
- ④ Learning from previous campaigns not straightforward
  - ▶ What if I previously optimized  $f_1(x)$  alone, does that help in optimizing  $[f_1(x), f_2(x)]$ ?
  - ▶ What if I previously optimized  $f(x_1, \dots, x_d)$ , does that help in optimizing  $f(x_1, \dots, x_d, x_{d+1})$ ?

→ Use dimension-agnostic task embeddings

## Caveats

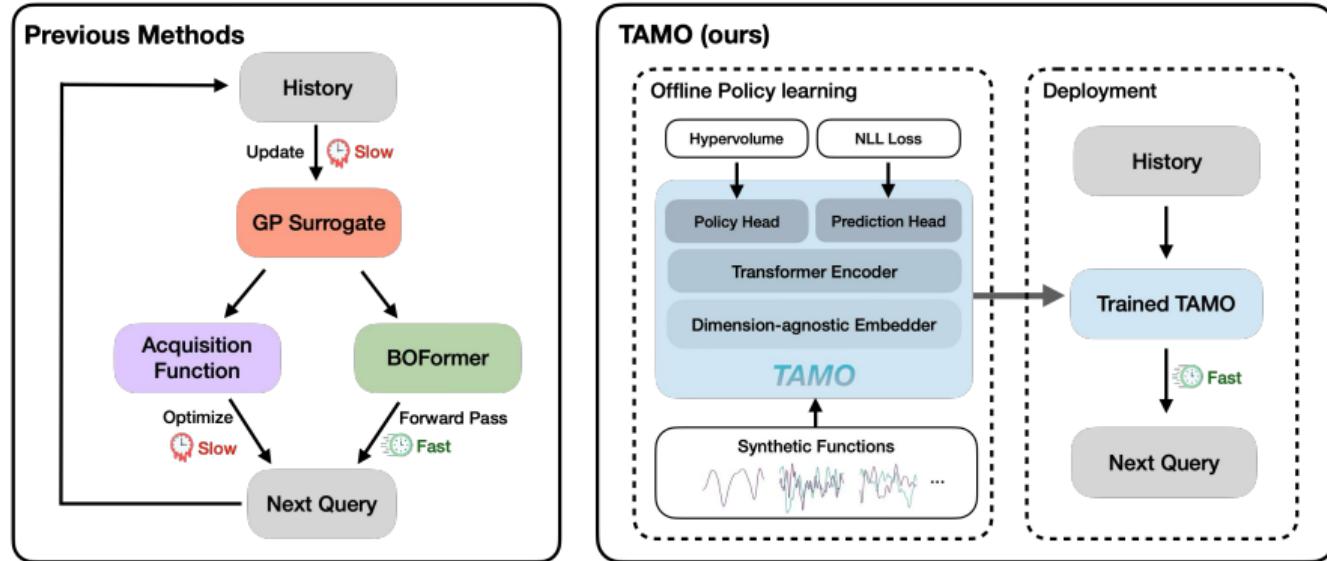
- ➊ Performances highly dependent on surrogate and acquisition function pair
  - ▶ Requires careful, expert selection of the ideal combination
  - ▶ Will not transfer to the next problem

→ "Learn" the language of optimization
- ➋ Slow in high-throughput settings due to vanilla GP cubic complexity  
→ Do a huge offline pre-training step, reduce inference to a feedforward pass
- ➌ Most of the time myopic, focused on 1-step optimality  
→ Train over long optimization horizons using RL
- ➍ Learning from previous campaigns not straightforward
  - ▶ What if I previously optimized  $f_1(x)$  alone, does that help in optimizing  $[f_1(x), f_2(x)]$ ?
  - ▶ What if I previously optimized  $f(x_1, \dots, x_d)$ , does that help in optimizing  $f(x_1, \dots, x_d, x_{d+1})$ ?

→ Use dimension-agnostic task embeddings

Solution: Task-Agnostic Amortization 😱

# Overview



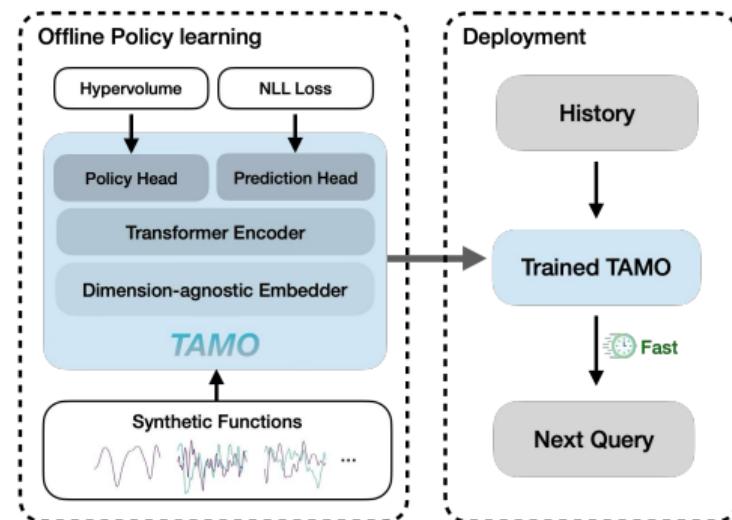
	Multi-objective	End-to-end amortized	Input agnostic	Output agnostic
Vanilla MOBO (Daulton et al., 2020)	✓	✗	N/A	N/A
BOFormer (Hung et al., 2025)	✓	✗	N/A	✗
NAP (Maraval et al., 2023)	✗	✓	✗	✗
DANP (Lee et al., 2025)	✗	✗	✓	✗
<b>TAMO (this work)</b>	✓	✓	✓	✓

# TAMO: architecture at a glance

Shared backbone. Each forward pass uses either a *prediction batch* (*context, targets*) or an *optimization batch* (*history, queries*), traversing the *same core components*.

**Optimization**  $\Leftrightarrow$  Predicting the optimum.

- ① **Dimension-agnostic embedder:** scalar $\rightarrow$ vector maps producing tokens independent of input/output dim.
- ② **Transformer encoder:** aggregates variable-size histories/contexts into a compact summary.
- ③ **Task conditioning:** a few learned tokens injected late to specialize the computation.
- ④ **Two heads:** a prediction head (density) and a policy head (acquisition over query set).



# Pretraining

**Task distribution.** Synthetic MOO tasks

$$\tau \sim p(\tau) \text{ with } f_\tau : \mathcal{X} \subset \mathbb{R}^{d_x^\tau} \rightarrow \mathbb{R}^{d_y^\tau}$$

Heterogeneous  $d_x^\tau, d_y^\tau \Rightarrow$   
dimension-agnostic policy.

Per step: two distinct mini-batches

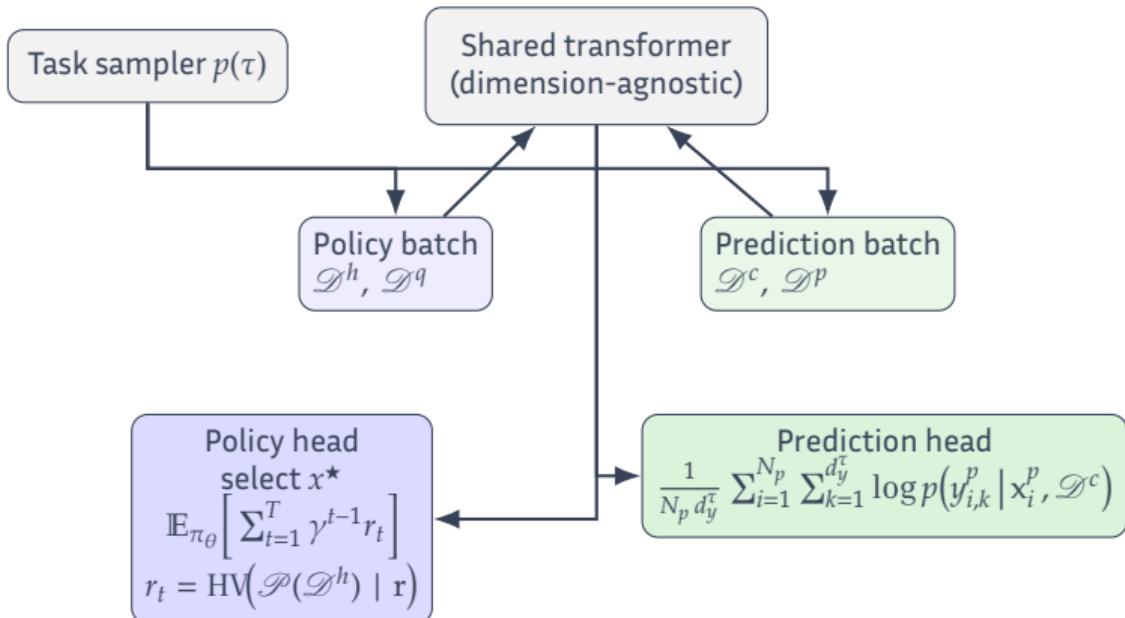
- *Policy-learning:*

$$\text{history } \mathcal{D}^h = \{(x^h, y^h)\}_{h=1}^{N_h}$$

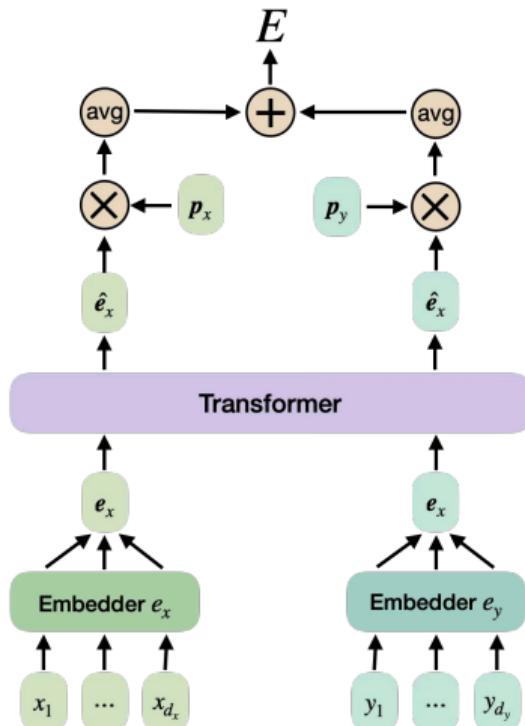
$$\text{query set } \mathcal{D}^q = \{x^q\}_{q=1}^{N_q}$$

policy picks the best  $x^q$  given  $\mathcal{D}^h$ .

- *Prediction:* fresh function draw;  
sample  $N$  pairs and split into  
context  $\mathcal{D}^c = \{(x^c, y^c)\}_{c=1}^{N_c}$  and  
targets  $\mathcal{D}^p = \{x^p\}_{p=1}^{N_p}$   
for in-context regression.



# Dimension-agnostic embedder



- Learnable maps (NNs)  $e_x, e_y : \mathbb{R} \rightarrow \mathbb{R}^{d_e}$  applied element-wise:  
 $e_x = e_x(x) \in \mathbb{R}^{d_x^\tau \times d_e}, \quad e_y = e_y(y) \in \mathbb{R}^{d_y^\tau \times d_e}.$
- After  $L$  transformer layers over  $[e_x; e_y]$  we get contextualized tokens  $\hat{e}_x, \hat{e}_y$ . Draw per-dimension tokens  $p_x^{(j)}, p_y^{(k)} \in \mathbb{R}^{d_e}$  sampled per batch from fixed pools of learned embeddings, then:

$$\tilde{e}_{x,j} = \hat{e}_{x,j} \odot p_x^{(j)}, \quad \tilde{e}_{y,k} = \hat{e}_{y,k} \odot p_y^{(k)}$$

$$\bar{e}_x = \frac{1}{d_x^\tau} \sum_{j=1}^{d_x^\tau} \tilde{e}_{x,j}, \quad \bar{e}_y = \frac{1}{d_y^\tau} \sum_{k=1}^{d_y^\tau} \tilde{e}_{y,k}, \quad E = \bar{e}_x + \bar{e}_y.$$

→ breaks permutation symmetry; features/objectives remain identifiable.

# Transformer encoder-decoder

Transformer layers split into  $B = B_1 + B_2$

- $B_1$ : history/context tokens self-attend  $\Rightarrow \hat{E}^h$  or  $\hat{E}^c$ ; queries/targets cross-attend to them  $\Rightarrow \hat{E}^q$  or  $\hat{E}^p$ .  
This is the *only* path for queries/targets to use past data.
- $B_2$ : drop history/context; keep query/target tokens + *task-specific tokens*. An attention mask enforces that query/target tokens only attend to task-specific tokens.

## Task-specific tokens

- *Prediction*: a prediction-task token and an output-index token  $p_y^{(k)}$ .
- *Optimization*: an optimization-task token, a time-budget token  $g_{\text{time}} = \text{MLP}_\theta((T-t)/T)$ , and an input-dimension token  $\sum_{j=1}^{d_x^T} p_x^{(j)}$ .

## Heads

*Prediction head (per scalar target):*

$$\{\phi_{i\ell}, \mu_{i\ell}, \sigma_{i\ell}\}_{\ell=1}^K = \text{MLP}_\theta(\hat{E}_i^p)$$

$$p(y_{i,k}^p \mid x_i^p, \mathcal{D}^c) = \sum_{\ell=1}^K \phi_{i\ell} \mathcal{N}(y_{i,k}^p; \mu_{i\ell}, \sigma_{i\ell}^2)$$

*Policy head (over queries):*

$$\alpha_i = \text{MLP}_\theta(\hat{E}_i^q)$$

$$\pi_\theta(x_i^q) = \frac{e^{\alpha_i}}{\sum_r e^{\alpha_r}}$$

---

**Algorithm S2** TAMO Test-Time Algorithm

---

**Require:** Pre-trained TAMO model, new task  $\tau_{\text{test}}$ , query budget  $T$ , initial history set  $\mathcal{D}_0^h := \{\mathbf{x}^h, y^h\}$  (with random samples if empty),

- 1:  $\mathcal{D}^h \leftarrow \mathcal{D}_0^h$  ▷ Initialize the history set
  - 2:  $\mathcal{P} \leftarrow \{y^h\}$  ▷ Initialize the Pareto set
  - 3: **for**  $t = 1, \dots, T$  **do**
  - 4:    $\mathbf{x}_t \sim \pi_\theta(\cdot \mid \mathcal{D}^h, t, T)$  ▷ Sample the next query location
  - 5:    $y_t \leftarrow \text{Evaluate}(\mathbf{x}_t, \tau_{\text{test}})$
  - 6:    $\mathcal{D}^h \leftarrow \mathcal{D}^h \cup \{(\mathbf{x}_t, y_t)\}$  ▷ Update the history set
  - 7:    $\mathcal{P} \leftarrow \mathcal{P} \cup \{y_t\}$  ▷ Update the Pareto set with the new observation
  - 8: **end for**
  - 9: **return**  $\mathcal{D}^h, \mathcal{P}$
-

## Pre-training dataset composition

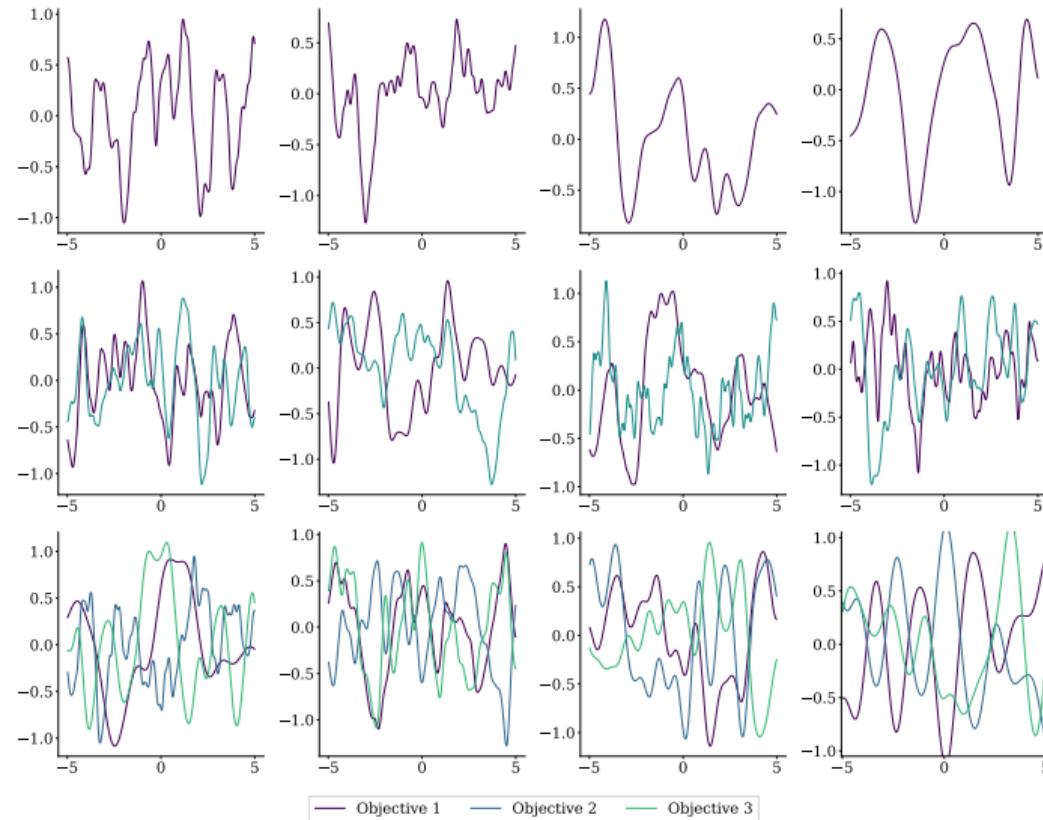
- Input dimensionality  $d_x \sim \mathcal{U}(\{1, 2\})$  and output dimensionality  $d_y \sim \mathcal{U}(\{1, 2, 3\})$ .
- For output correlations, with  $p = 1/2$ , either:
  - independent output dimensions are sampled
  - drawn from a multi-task GP, with task covariance defined as  $k(i, j) = (\mathbf{C}\mathbf{C}^T + \text{diag}(\mathbf{v}))_{i,j}$ ,  $i, j \in \{1, \dots, d_y\}$ , with  $\mathbf{C}$  is a low-rank matrix with rank  $r \sim \mathcal{U}(\{1, \dots, d_y\})$ .
- Data kernel along each output dimension:
  - Equally sampled from RBF, Matérn-3/2, Matérn-5/2
  - Standard deviation  $\sigma \sim U([0.1, 1.0])$
  - Lengthscale  $\ell \sim \mathcal{N}(2/3, 0.5)$  truncated to  $[0.1, 2.0]$ .
- The sampled function values  $\mathbf{y}$  were centered and normalized to lie within  $[-1, 1]^{d_y}$ .

## Pre-training dataset composition

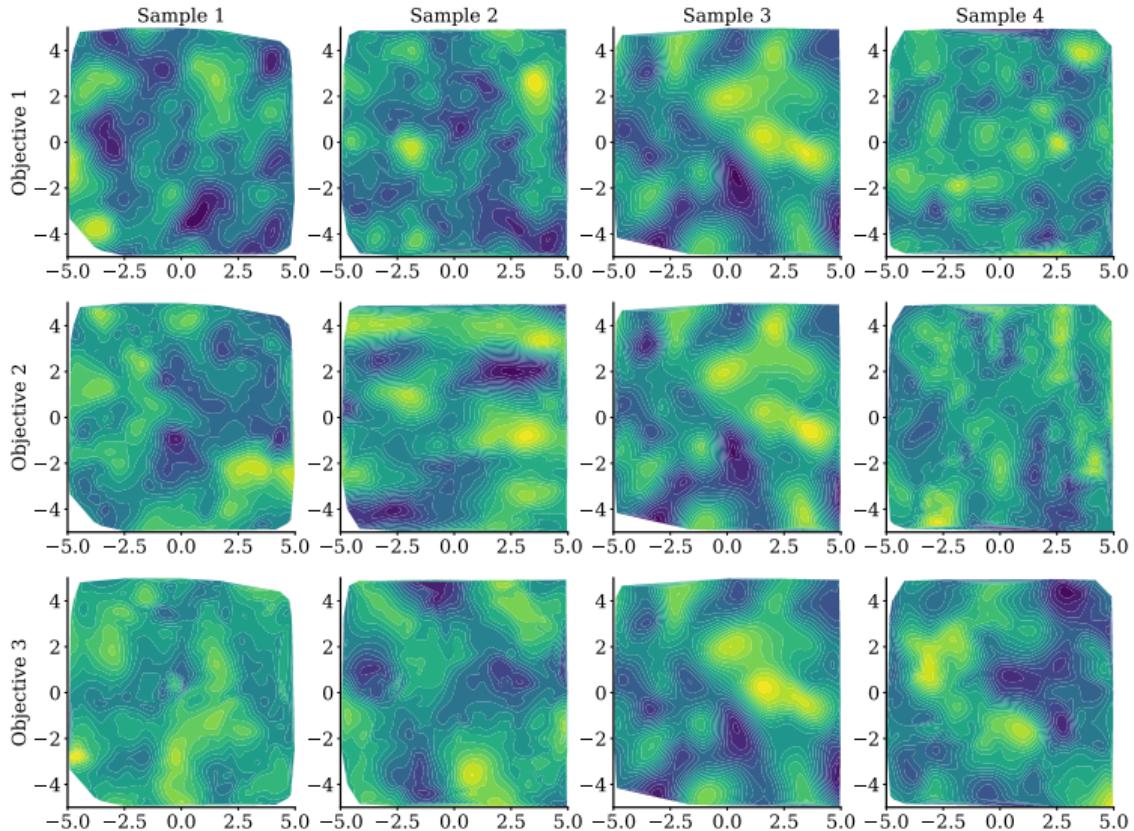
- Input dimensionality  $d_x \sim \mathcal{U}(\{1, 2\})$  and output dimensionality  $d_y \sim \mathcal{U}(\{1, 2, 3\})$ .
- For output correlations, with  $p = 1/2$ , either:
  - ▶ independent output dimensions are sampled
  - ▶ drawn from a multi-task GP, with task covariance defined as  $k(i, j) = (\mathbf{C}\mathbf{C}^T + \text{diag}(\mathbf{v}))_{i,j}$ ,  $i, j \in \{1, \dots, d_y\}$ , with  $\mathbf{C}$  is a low-rank matrix with rank  $r \sim \mathcal{U}(\{1, \dots, d_y\})$ .
- Data kernel along each output dimension:
  - ▶ Equally sampled from RBF, Matérn-3/2, Matérn-5/2
  - ▶ Standard deviation  $\sigma \sim U([0.1, 1.0])$
  - ▶ Lengthscale  $\ell \sim \mathcal{N}(2/3, 0.5)$  truncated to  $[0.1, 2.0]$ .
- The sampled function values  $\mathbf{y}$  were centered and normalized to lie within  $[-1, 1]^{d_y}$ .

Completely synthetic dataset! 😱

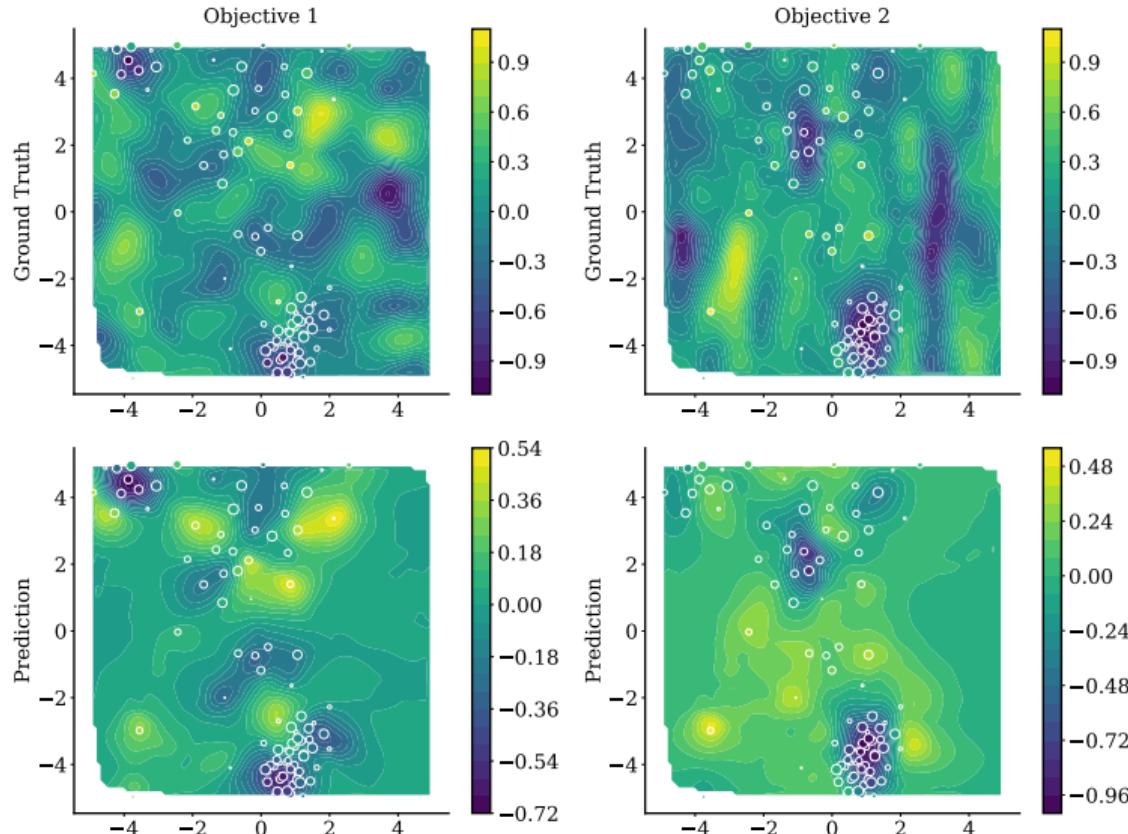
## Samples from pre-training dataset, $d_x = 1, d_y = 3$



# Samples from pre-training dataset, $d_x = 2, d_y = 3$



# Optimization run example, $d_x = 2, d_y = 2$



# Behind the scenes



Dimension-Agnostic Embedder	
Number of learnable positional tokens for $x$	4
Number of learnable positional tokens for $y$	3
Number of Transformer layers ( $L$ )	4
Dimension of $\mathbf{e}_x$ and $\mathbf{e}_y$	64
Transformer Encoder-Decoder	
Dimension of Transformer inputs	64
Point-wise feed-forward dimension of Transformer	256
Number of self-attention layers in Transformer ( $B$ )	8
Number of self-attention heads in Transformer	4
Heads	
Number of hidden layers in policy head	3
Number of components in GMM head ( $K$ )	20
Number of hidden layers in MLP for each GMM component	3
Training	
Number of iterations	400000
Number of burn-in iterations	393500
Initial learning rate for warm-up iterations ( $\text{lr}_1$ )	$1 \cdot 10^{-4}$
Initial Learning rate after warm-up ( $\text{lr}_2$ )	$4 \cdot 10^{-5}$
Learning rate scheduling	Linearly increase from 0 to $\text{lr}_1$ in the first 5% of total iterations; Cosine decay to 0 over total iterations
Size of prediction batch	32
Size of policy-learning batch	16
Weight on prediction loss ( $\lambda_{\text{rl}}$ )	1.0
discount factor ( $\gamma$ )	1.0
Size of context set	$N_c \sim U[2, 50 \cdot d_x^{\tau}]$
Size of target set ( $N_t$ )	$300 - N_c$
Size of query set ( $N_q$ )	256
Optimization budget $T$	100
Noise level $\sigma$	0.0
Number of initial observations during pretraining	1
Evaluation	
Number of initial observations during test time	1
Noise level $\sigma$	0.0
Size of query set ( $N_q$ )	2048
Optimization budget ( $T$ )	100

- Positional tokens define max dimensionality the model can handle

# Behind the scenes



Dimension-Agnostic Embedder	
Number of learnable positional tokens for $x$	4
Number of learnable positional tokens for $y$	3
Number of Transformer layers ( $L$ )	4
Dimension of $\mathbf{e}_x$ and $\mathbf{e}_y$	64
Transformer Encoder-Decoder	
Dimension of Transformer inputs	64
Point-wise feed-forward dimension of Transformer	256
Number of self-attention layers in Transformer ( $B$ )	8
Number of self-attention heads in Transformer	4
Heads	
Number of hidden layers in policy head	3
Number of components in GMM head ( $K$ )	20
Number of hidden layers in MLP for each GMM component	3
Training	
Number of iterations	400000
Number of burn-in iterations	393500
Initial learning rate for warm-up iterations ( $\text{lr}_1$ )	$1 \cdot 10^{-4}$
Initial Learning rate after warm-up ( $\text{lr}_2$ )	$4 \cdot 10^{-5}$
Learning rate scheduling	Linearly increase from 0 to $\text{lr}_1$ in the first 5% of total iterations; Cosine decay to 0 over total iterations
Size of prediction batch	32
Size of policy-learning batch	16
Weight on prediction loss ( $\lambda_{\text{rl}}$ )	1.0
discount factor ( $\gamma$ )	1.0
Size of context set	$N_c \sim U[2, 50 \cdot d_x^{\tau}]$
Size of target set ( $N_t$ )	$300 - N_c$
Size of query set ( $N_q$ )	256
Optimization budget $T$	100
Noise level $\sigma$	0.0
Number of initial observations during pretraining	1
Evaluation	
Number of initial observations during test time	1
Noise level $\sigma$	0.0
Size of query set ( $N_q$ )	2048
Optimization budget ( $T$ )	100

- Positional tokens define max dimensionality the model can handle
- For comparison: TabPFN uses  $L = 12$ ,  $\mathbf{e}_x \in \mathbb{R}^{512}$  to handle  $d_x \approx 50$

# Behind the scenes



Dimension-Agnostic Embedder	
Number of learnable positional tokens for $x$	4
Number of learnable positional tokens for $y$	3
Number of Transformer layers ( $L$ )	4
Dimension of $\mathbf{e}_x$ and $\mathbf{e}_y$	64
Transformer Encoder-Decoder	
Dimension of Transformer inputs	64
Point-wise feed-forward dimension of Transformer	256
Number of self-attention layers in Transformer ( $B$ )	8
Number of self-attention heads in Transformer	4
Heads	
Number of hidden layers in policy head	3
Number of components in GMM head ( $K$ )	20
Number of hidden layers in MLP for each GMM component	3
Training	
Number of iterations	400000
Number of burn-in iterations	393500
Initial learning rate for warm-up iterations ( $\text{lr}_1$ )	$1 \cdot 10^{-4}$
Initial Learning rate after warm-up ( $\text{lr}_2$ )	$4 \cdot 10^{-5}$
Learning rate scheduling	Linearly increase from 0 to $\text{lr}_1$ in the first 5% of total iterations; Cosine decay to 0 over total iterations
Size of prediction batch	32
Size of policy-learning batch	16
Weight on prediction loss ( $\lambda_{\text{rl}}$ )	1.0
discount factor ( $\gamma$ )	1.0
Size of context set	$N_c \sim U[2, 50 \cdot d_x^{\tau}]$
Size of target set ( $N_t$ )	$300 - N_c$
Size of query set ( $N_q$ )	256
Optimization budget $T$	100
Noise level $\sigma$	0.0
Number of initial observations during pretraining	1
Evaluation	
Number of initial observations during test time	1
Noise level $\sigma$	0.0
Size of query set ( $N_q$ )	2048
Optimization budget ( $T$ )	100

- Positional tokens define max dimensionality the model can handle
- For comparison: TabPFN uses  $L = 12$ ,  $\mathbf{e}_x \in \mathbb{R}^{512}$  to handle  $d_x \approx 50$
- $\approx 1.1\text{M}$  parameters

# Behind the scenes



Dimension-Agnostic Embedder	
Number of learnable positional tokens for $x$	4
Number of learnable positional tokens for $y$	3
Number of Transformer layers ( $L$ )	4
Dimension of $\mathbf{e}_x$ and $\mathbf{e}_y$	64
Transformer Encoder-Decoder	
Dimension of Transformer inputs	64
Point-wise feed-forward dimension of Transformer	256
Number of self-attention layers in Transformer ( $B$ )	8
Number of self-attention heads in Transformer	4
Heads	
Number of hidden layers in policy head	3
Number of components in GMM head ( $K$ )	20
Number of hidden layers in MLP for each GMM component	3
Training	
Number of iterations	400000
Number of burn-in iterations	393500
Initial learning rate for warm-up iterations ( $\text{lr}_1$ )	$1 \cdot 10^{-4}$
Initial Learning rate after warm-up ( $\text{lr}_2$ )	$4 \cdot 10^{-5}$
Learning rate scheduling	Linearly increase from 0 to $\text{lr}_1$ in the first 5% of total iterations; Cosine decay to 0 over total iterations
Size of prediction batch	32
Size of policy-learning batch	16
Weight on prediction loss ( $\lambda_{\text{rl}}$ )	1.0
discount factor ( $\gamma$ )	1.0
Size of context set	$N_c \sim U[2, 50 \cdot d_x^{\tau}]$
Size of target set ( $N_t$ )	$300 - N_c$
Size of query set ( $N_q$ )	256
Optimization budget $T$	100
Noise level $\sigma$	0.0
Number of initial observations during pretraining	1
Evaluation	
Number of initial observations during test time	1
Noise level $\sigma$	0.0
Size of query set ( $N_q$ )	2048
Optimization budget ( $T$ )	100

- Positional tokens define max dimensionality the model can handle
- For comparison: TabPFN uses  $L = 12$ ,  $\mathbf{e}_x \in \mathbb{R}^{512}$  to handle  $d_x \approx 50$
- $\approx 1.1\text{M}$  parameters
- Training time  $\approx 2$  days
- The model “learns” from 100-iterations tasks  $\implies$  outside is OOD!

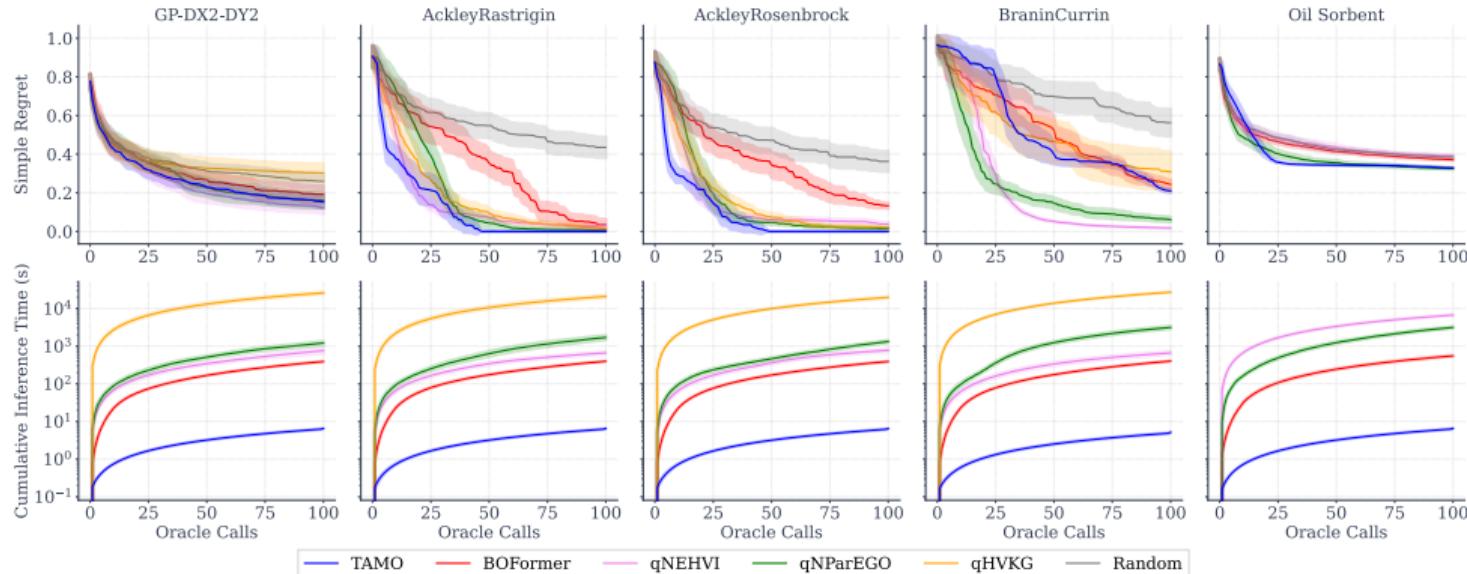
# Behind the scenes



Dimension-Agnostic Embedder	
Number of learnable positional tokens for $x$	4
Number of learnable positional tokens for $y$	3
Number of Transformer layers ( $L$ )	4
Dimension of $\mathbf{e}_x$ and $\mathbf{e}_y$	64
Transformer Encoder-Decoder	
Dimension of Transformer inputs	64
Point-wise feed-forward dimension of Transformer	256
Number of self-attention layers in Transformer ( $B$ )	8
Number of self-attention heads in Transformer	4
Heads	
Number of hidden layers in policy head	3
Number of components in GMM head ( $K$ )	20
Number of hidden layers in MLP for each GMM component	3
Training	
Number of iterations	400000
Number of burn-in iterations	393500
Initial learning rate for warm-up iterations ( $\text{lr}_1$ )	$1 \cdot 10^{-4}$
Initial Learning rate after warm-up ( $\text{lr}_2$ )	$4 \cdot 10^{-5}$
Learning rate scheduling	Linearly increase from 0 to $\text{lr}_1$ in the first 5% of total iterations; Cosine decay to 0 over total iterations
Size of prediction batch	32
Size of policy-learning batch	16
Weight on prediction loss ( $\lambda_{\text{rl}}$ )	1.0
discount factor ( $\gamma$ )	1.0
Size of context set	$N_c \sim U[2, 50 \cdot d_x^{\tau}]$
Size of target set ( $N_t$ )	$300 - N_c$
Size of query set ( $N_q$ )	256
Optimization budget $T$	100
Noise level $\sigma$	0.0
Number of initial observations during pretraining	1
Evaluation	
Number of initial observations during test time	1
Noise level $\sigma$	0.0
Size of query set ( $N_q$ )	2048
Optimization budget ( $T$ )	100

- Positional tokens define max dimensionality the model can handle
- For comparison: TabPFN uses  $L = 12$ ,  $\mathbf{e}_x \in \mathbb{R}^{512}$  to handle  $d_x \approx 50$
- $\approx 1.1\text{M}$  parameters
- Training time  $\approx 2$  days
- The model “learns” from 100-iterations tasks  $\implies$  outside is OOD!
- Query set = grid for evaluations!  
Increases inference time linearly.

# Results



**Figure 3: Synthetic and real-world multi-objective benchmarks:** simple regret (top) and cumulative inference time (bottom) vs. oracle calls (mean  $\pm$  95% CIs over 30 runs). **TAMO achieves competitive regret while cutting proposal time by  $50\times$ – $1000\times$ .**

## Results - Single-objective BO

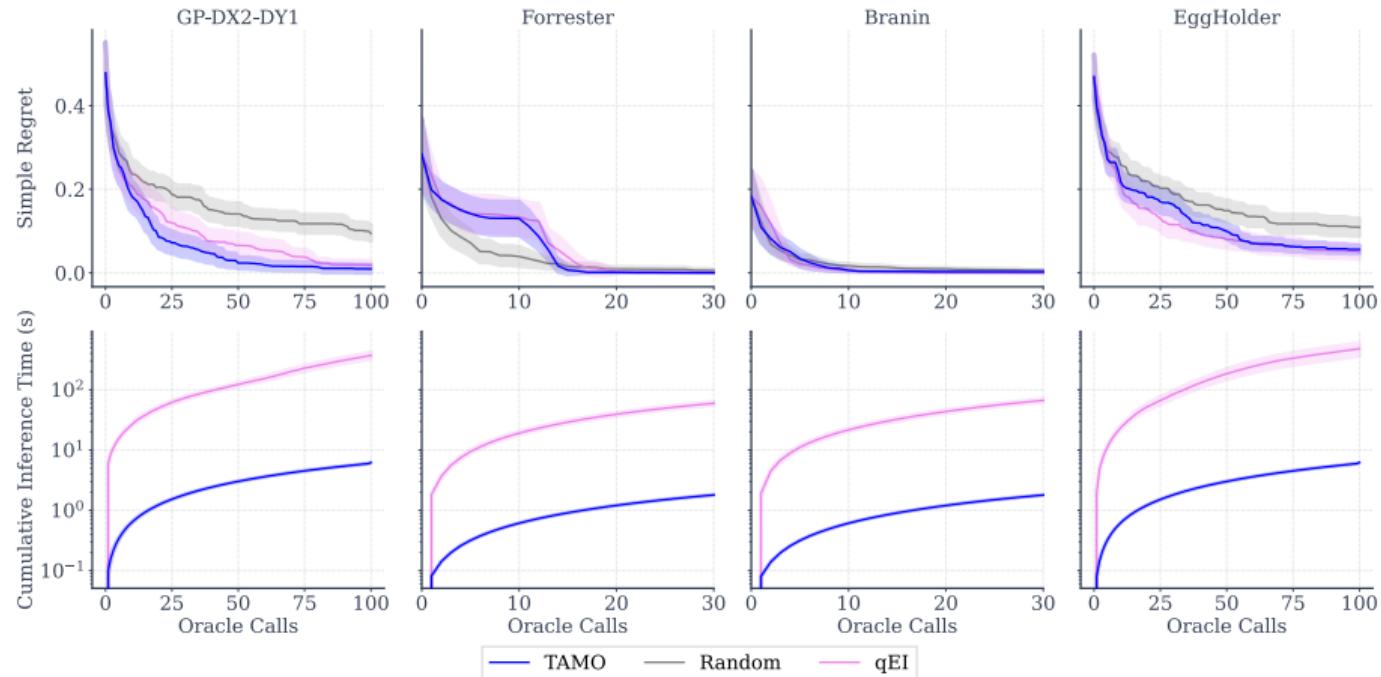


Figure S1: Simple regret and inference time on **synthetic examples for single-objective optimization**. Mean with 95% confidence intervals computed across 30 runs with random initial observations. Again, TAMO matches state-of-the-art regret while dramatically reducing proposal time.

# Results - Out of distribution examples

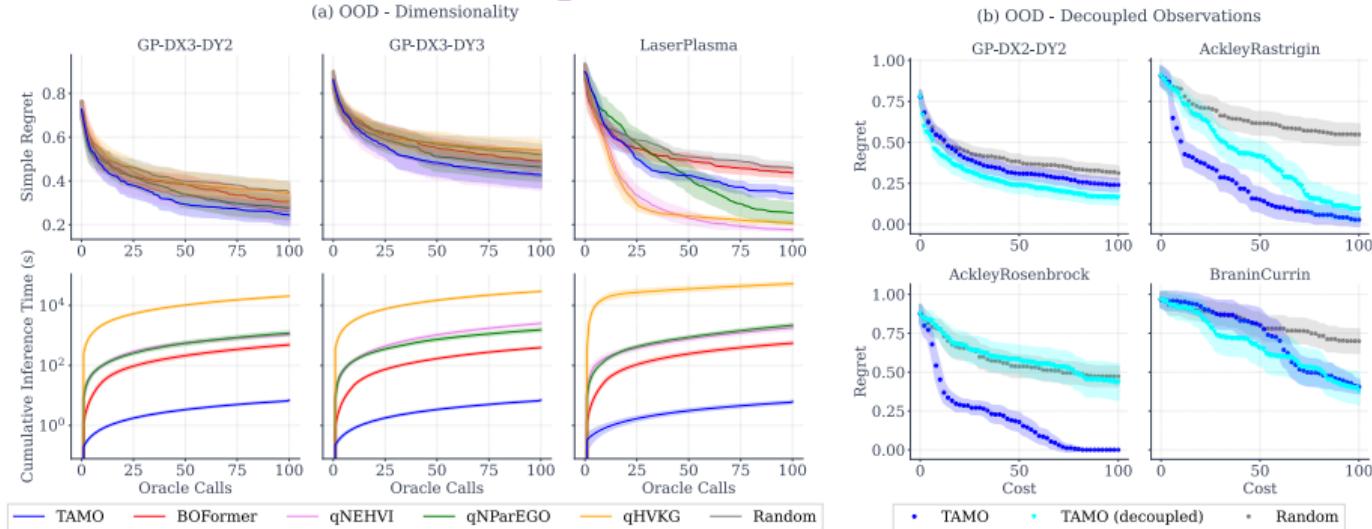


Figure 4: **Out-of-distribution evaluations.** **(a) Dimensionality:** simple regret (top) and cumulative inference time (bottom) on tasks whose input/output dimensions are unseen at pretraining. **(b) Decoupled observations:** regret vs. *cumulative cost* when, at step  $t$ , the optimizer may observe both objectives at cost 2 (dark blue) or only one at cost 1 (cyan). Curves show means with 95% confidence intervals over 30 runs with random initial observations. **TAMO offers promising generalization capabilities across unseen dimensionalities and decoupled feedback settings, delivering orders-of-magnitude faster proposals while maintaining competitive regret.**

## Conclusion, Limitations, Perspectives

### Universal black-box optimization model

- Extending to constrained, multi-source, preferential, cost-aware, etc... Does not require a paradigm shift, mostly pre-training modification
- Sky's the limit: just make the model bigger, increase pretraining dataset size, train longer

# Conclusion, Limitations, Perspectives

## Universal black-box optimization model

- Extending to constrained, multi-source, preferential, cost-aware, etc... Does not require a paradigm shift, mostly pre-training modification
- Sky's the limit: just make the model bigger, increase pretraining dataset size, train longer

Actually...

- The policy will quickly overfit 😱

# Conclusion, Limitations, Perspectives

## Universal black-box optimization model

- Extending to constrained, multi-source, preferential, cost-aware, etc... Does not require a paradigm shift, mostly pre-training modification
- Sky's the limit: just make the model bigger, increase pretraining dataset size, train longer

Actually...

- The policy will quickly overfit 😬
- Does not scale yet 🤯

# Conclusion, Limitations, Perspectives

## Universal black-box optimization model

- Extending to constrained, multi-source, preferential, cost-aware, etc... Does not require a paradigm shift, mostly pre-training modification
- Sky's the limit: just make the model bigger, increase pretraining dataset size, train longer

Actually...

- The policy will quickly overfit 😬
- Does not scale yet 🤯
- Highly black box method 😐

# Conclusion, Limitations, Perspectives

## Universal black-box optimization model

- Extending to constrained, multi-source, preferential, cost-aware, etc... Does not require a paradigm shift, mostly pre-training modification
- Sky's the limit: just make the model bigger, increase pretraining dataset size, train longer

Actually...

- The policy will quickly overfit 😬
- Does not scale yet 🤯
- Highly black box method 😐

Enough work avenues to fill several PhD theses

- Scaling to high-dimensional settings, handling structured objects
- Explainability
- Pre-training dataset  $\iff$  Prior from a Bayesian perspective. Its composition is key.

# Appendix



## Two-steps training

- ① Warm up backbone on prediction task by minimizing a negative log-likelihood over  $(\mathcal{D}^c, \mathcal{D}^p)$

$$\mathcal{L}^{(p)}(\theta) = -\mathbb{E}_{\tau \sim p(\tau)} \left[ \frac{1}{N_p d_y^\tau} \sum_{i=1}^{N_p} \sum_{k=1}^{d_y^\tau} \log p(y_{i,k}^p | x_i^p, \mathcal{D}^c) \right]$$

→ promotes accurate in-context regression and useful representations

- ② Then, optimize the policy  $\pi_\theta(x | s)$  with trajectory objective

$$J(\theta) = \mathbb{E}_{\tau \sim p(\tau)} \left[ \mathbb{E}_{\pi_\theta} \left[ \sum_{t=1}^T \gamma^{t-1} r_t \right] \right]$$

→ aligns learning signal with improvements in Pareto quality alongside the prediction objective.

The overall objective combines both terms:

$$\mathcal{L}(\theta) = \mathcal{L}^{(p)}(\theta) + \lambda_{rl} \mathcal{L}^{(rl)}(\theta), \quad \mathcal{L}^{(rl)}(\theta) = -J(\theta),$$

- $\mathcal{L}^{(p)}$  and  $\mathcal{L}^{(rl)}$  calculated from two distinct forward passes with different datasets
- Then summed for one single backward pass
- Training on full trajectories directly rewards long-horizon improvements,
- Amortization enables learning from many tasks offline.

**Algorithm S1** TAMO Pre-Training Algorithm

---

**Require:** task distribution  $p(\tau)$ , prediction context size  $N_c$ , prediction target size  $N_p$ , query budget  $T$ , number of burn-in iterations  $\eta$ , number of total iterations num\_total\_iterations

```

1: for iteration  $i = 1, \dots, \text{num\_total\_iterations}$  do
2:   Sample a task  $\tau \sim p(\tau)$ 
3:   Sample prediction batches  $\mathcal{D}^c = \{(\mathbf{x}_i^c, y_i^c)\}_{i=1}^{N_c}$  and  $\mathcal{D}^p = \{\mathbf{x}_i^p\}_{i=1}^{N_p}$  from  $\tau$ 
4:   Model predicts outcomes:  $p(y_{i,k}^p \mid \mathbf{x}_i^p, \mathcal{D}^c), \forall \mathbf{x}_i^p \in \mathcal{D}^p$ 
5:   if  $i \leq \eta$  then
6:     Update model by minimizing the prediction loss  $\mathcal{L}^{(p)}$  (Equation 5)
7:   else
8:     Sample a new task  $\tau \sim p(\tau)$ 
9:     Sample query set  $\mathcal{D}^q$ 
10:    Initialize a history set  $\mathcal{D}^h \leftarrow \{(\mathbf{x}_0^h, y_0^h)\}, \mathbf{x}_0^h \sim \mathcal{D}^q$ 
11:    Set reference point  $\mathbf{r}$  and calculate optimal Hypervolume:  $\text{HV}^* \leftarrow \text{HV}(\mathcal{P}(\mathcal{X}) \mid \mathbf{r})$ 
12:    Initialize Pareto set  $\mathcal{P} \leftarrow \{y_0^h\}$ 
13:    for  $t = 1, \dots, T$  do
14:      Select next query point:  $\mathbf{x}_t \sim \pi_\theta(\cdot \mid \mathcal{D}^h, t, T)$ 
15:       $y_t \leftarrow \text{Evaluate}(\mathbf{x}_t, \tau)$ 
16:      Update history set:  $\mathcal{D}^h \leftarrow \mathcal{D}^h \cup \{(\mathbf{x}_t, y_t)\}$ 
17:      Update Pareto set:  $\mathcal{P} \leftarrow \mathcal{P} \cup \{y_t\}$ 
18:      Compute reward:  $r_t = \frac{\text{HV}(\mathcal{P} \mid \mathbf{r})}{\text{HV}^*}$ 
19:    end for
20:    Update model using the overall objective  $\mathcal{L}$  (Equation 6)
21:  end if
22: end for

```

---